

# Final Project Milestone 3

April 15, 2022

- CIS 9440 - Data Warehousing for Analytics
- Final Project Milestone 3
- Group Number - **17**
- Student – **KWANG HEUM YEON**

## 1 ETL Process

This ETL process is for the final project of CIS 9440, Zicklin School of Business in Baruch College. The overall guideline and most of the codes for the process was provided by the professor. The KPI of the project is as stated in project milestone #2. There are three separate datasets to get the KPI. Two of them were extracted from the NYC open data(<https://opendata.cityofnewyork.us/>) and the other referred to New York Demographics by Cubit ([https://www.newyork-demographics.com/zip\\_codes\\_by\\_population](https://www.newyork-demographics.com/zip_codes_by_population)). We have 10 stages to get the planned result and those are briefly as follows:

- (1) Extract the first two datasets, merge, and clean
- (2) Merge the third dataset and clean (\*We will combine the third dataset with the completed set narrowed down the range of the dataset to avoid file size issues.)
- (3) Create fact tables and each dimension tables
- (4) Deliver fact tables and each dimension table to the data warehouse (Big Query)
- (5) Appendix

Each code cell will contain a comment. In the appendix, we will attach the test images screen-captured from the data warehouse (Big Query) and BI application (Tableau) to check if the delivered tables are properly recognized.

### 1.1 Step 1: Extract data

1. connect to NYC Open Data with API Key
2. pull specific dataset as a pandas dataframe
3. Look at shape of extracted data

```
[1]: # import libraries
import pandas as pd
from sodapy import Socrata
from google.cloud import bigquery
```

```
from google.oauth2 import service_account
```

### 1.1.1 Bicycle Counts

```
[2]: data_url = 'data.cityofnewyork.us'    # The Host Name for the API endpoint (the
      ↪https:// part will be added automatically)
      data_set = 'uczf-rk3c'    # The data set at the API endpoint
      app_token = 'EBzL4gV7ZLeOF6y5zsvwWYm3M'    # The App Token code created in the
      ↪prior steps

      # full URL to look at data on NYC Open Data
      # https://data.cityofnewyork.us/resource/uczf-rk3c.json
```

```
[3]: # create the client that points to the API endpoint
      client = Socrata(data_url, app_token, timeout = 200)    #time limit run this cell:
      ↪ 200 sec
```

```
[4]: print(f"client name is: {client}")
      print(f"client data type is: {type(client)}")
```

client name is: <sodapy.socrata.Socrata object at 0x000001CB2A3D4130>  
client data type is: <class 'sodapy.socrata.Socrata'>

```
[5]: # test the connection to NYC Open Data

      # retrieve the first 100 rows from the data_set
      test_results = client.get(data_set, limit = 100)

      # the test_results are returned as JSON object from the API
      # the sodapy library converts this JSON object to a python list of dictionaries
      # now, convert the list of dictionaries to a pandas data frame
      test_results_df = pd.DataFrame.from_records(test_results)
```

```
[6]: # examine the test_results_df pandas dataframe
      test_results_df.head()
```

```
[6]:   id1 counts      date status      site
0    0     41  2012-08-31T00:00:00.000      4  100005020
1    1     52  2012-08-31T00:15:00.000      4  100005020
2    2     38  2012-08-31T00:30:00.000      4  100005020
3    3     36  2012-08-31T00:45:00.000      4  100005020
4    4     40  2012-08-31T01:00:00.000      4  100005020
```

sodapy client.get parameters 1. select 2. where 3. order 4. limit 5. group

```
[7]: # next, get the total number of records in our the entire data set
      total_record_count = client.get(data_set, select = "COUNT(*)")
      print(f"total records in {data_set}: {total_record_count}")
```

total records in ucfz-rk3c: [{'COUNT': '4628092'}]

```
[8]: # next, get the total number of records in our target data set
target_record_count = client.get(data_set,
                                where = "date > '2021-01-01'",
                                select= "COUNT(*)")
print(f"target records in {data_set}: {target_record_count}")
```

target records in ucfz-rk3c: [{'COUNT': '620494'}]

```
[9]: # loop through data set to pull all rows in chunks (cannot pull all rows at
    ↳ once)

# measure time this function takes
import time
start_time = time.time()

start = 0          # start at 0
chunk_size = 2000  # fetch 2000 rows at a time
results = []       # empty out our result list
record_count = target_record_count

while True:

    # fetch the set of records starting at 'start'
    results.extend(client.get(data_set,
                              where = "date > '2021-01-01'",
                              offset = start,
                              limit = chunk_size))

    # update the starting record number
    start = start + chunk_size

    # if we have fetched all of the records (we have reached record_count),
    ↳ exit loop
    if (start > int(record_count[0]['COUNT'])):
        break

# convert the list into a pandas data frame
data = pd.DataFrame.from_records(results)

end_time = time.time()
print(f"loop to {round(end_time - start_time, 1)} seconds")

data.info()
```

loop to 102.8 seconds

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 620494 entries, 0 to 620493

Data columns (total 5 columns):

#	Column	Non-Null	Count	Dtype
0	id1	620494	non-null	object
1	counts	620360	non-null	object
2	date	620494	non-null	object
3	status	620360	non-null	object
4	site	620494	non-null	object

dtypes: object(5)

memory usage: 23.7+ MB

```
[10]: data.head(2)
```

```
[10]:      id1 counts      date status      site
0  145537     11  2021-01-01T00:15:00.000      0  100009425
1   248161      5  2021-01-01T00:15:00.000      0  100009426
```

```
[11]: count = data.copy()
```

### 1.1.2 Bicycle Sites

```
[12]: data_url = 'data.cityofnewyork.us'      # The Host Name for the API endpoint (the
      ↪ https:// part will be added automatically)
data_set = 'smn3-rz9f9'      # The data set at the API endpoint
app_token = 'EBzL4gV7ZLe0F6y5zsvwWYm3M'    # The App Token code created in the
      ↪ prior steps

# full URL to look at data on NYC Open Data
# https://data.cityofnewyork.us/resource/smn3-rz9f9.json
```

```
[13]: # create the client that points to the API endpoint
client = Socrata(data_url, app_token, timeout = 200) #time limit run this cell:
      ↪ 200 sec
```

```
[14]: print(f"client name is: {client}")
      print(f"client data type is: {type(client)}")
```

client name is: <sodapy.socrata.Socrata object at 0x000001CB2A467790>

client data type is: <class 'sodapy.socrata.Socrata'>

```
[15]: # test the connection to NYC Open Data

# retrieve the first 100 rows from the data_set
test_results = client.get(data_set, limit = 100)

# the test_results are returned as JSON object from the API
# the sodapy library converts this JSON object to a python list of dictionaries
```

```
# now, convert the list of dictionaries to a pandas data frame
test_results_df = pd.DataFrame.from_records(test_results)
```

```
[16]: # examine the test_results_df pandas dataframe
test_results_df.head(2)
```

```
[16]:   id          name  latitude  longitude \
0  0  Manhattan Bridge 2012 Test Bike Counter   40.69981  -73.98589
1  5    Ed Koch Queensboro Bridge Shared Path   40.751038  -73.94082

      domain      site      timezone interval \
0  New York City DOT  100005020  (UTC-05:00) US/Eastern;DST      15
1  New York City DOT  100009428  (UTC-05:00) US/Eastern;DST      15

      :@computed_region_efsh_h5xi :@computed_region_f5dn_yrer \
0                                16865                        68
1                                16858                        53

      :@computed_region_yeji_bk3q :@computed_region_92fq_4b7q \
0                                2                        38
1                                3                        33

      :@computed_region_sbqj_enih      counter
0                                54      NaN
1                                66  Y2H19111445
```

sodapy client.get parameters 1. select 2. where 3. order 4. limit 5. group

```
[17]: # next, get the total number of records in our the entire data set
total_record_count = client.get(data_set, select = "COUNT(*)")
print(f"total records in {data_set}: {total_record_count}")
```

total records in smn3-rzf9: [{'COUNT': '26'}]

```
[18]: # extract the entire data set
target_record_count = client.get(data_set,
                                  select= "COUNT(*)")
print(f"target records in {data_set}: {target_record_count}")
```

target records in smn3-rzf9: [{'COUNT': '26'}]

```
[19]: # loop through data set to pull all rows in chunks (cannot pull all rows at_
      ↪once)

# measure time this function takes
import time
start_time = time.time()

start = 0          # start at 0
```

```

chunk_size = 2000      # fetch 2000 rows at a time
results = []           # empty out our result list
record_count = target_record_count

while True:

    # fetch the set of records starting at 'start'
    results.extend(client.get(data_set,
                              offset = start,
                              limit = chunk_size))

    # update the starting record number
    start = start + chunk_size

    # if we have fetched all of the records (we have reached record_count),
    → exit loop
    if (start > int(record_count[0]['COUNT'])):
        break

# convert the list into a pandas data frame
data = pd.DataFrame.from_records(results)

end_time = time.time()
print(f"loop to {round(end_time - start_time, 1)} seconds")

data.info()

```

loop to 0.2 seconds

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 26 entries, 0 to 25

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	id	26 non-null	object
1	name	26 non-null	object
2	latitude	26 non-null	object
3	longitude	26 non-null	object
4	domain	26 non-null	object
5	site	26 non-null	object
6	timezone	26 non-null	object
7	interval	26 non-null	object
8	:@computed_region_efsh_h5xi	22 non-null	object
9	:@computed_region_f5dn_yrer	23 non-null	object
10	:@computed_region_yeji_bk3q	23 non-null	object
11	:@computed_region_92fq_4b7q	23 non-null	object
12	:@computed_region_sbqj_enih	23 non-null	object
13	counter	16 non-null	object

dtypes: object(14)

memory usage: 3.0+ KB

```
[20]: data.head(2)
```

```
[20]:   id          name  latitude  longitude \
0  0  Manhattan Bridge 2012 Test Bike Counter  40.69981  -73.98589
1  5    Ed Koch Queensboro Bridge Shared Path  40.751038  -73.94082

      domain      site      timezone interval \
0  New York City DOT  100005020  (UTC-05:00) US/Eastern;DST      15
1  New York City DOT  100009428  (UTC-05:00) US/Eastern;DST      15

      :@computed_region_efsh_h5xi :@computed_region_f5dn_yrer \
0          16865          68
1          16858          53

      :@computed_region_yeji_bk3q :@computed_region_92fq_4b7q \
0          2          38
1          3          33

      :@computed_region_sbqj_enih      counter
0          54      NaN
1          66  Y2H19111445
```

```
[21]: site = data.copy()
```

## 1.2 Step 2: Data Profiling

1. Distinct values per column
2. Null values per column
3. Summary statistics per numeric column

```
[22]: # merge dataset, 'counts' and 'sites'
data = pd.merge(count, site, on = 'site', how = 'inner')
original_data = data.copy()
data.head(2)
```

```
[22]:   id1 counts      date status      site id \
0  145537    11  2021-01-01T00:15:00.000    0  100009425  2
1  145538     8  2021-01-01T00:30:00.000    0  100009425  2

      name      latitude      longitude      domain \
0  Prospect Park West  40.67128846  -73.97138165  New York City DOT
1  Prospect Park West  40.67128846  -73.97138165  New York City DOT

      timezone interval :@computed_region_efsh_h5xi \
0  (UTC-05:00) US/Eastern;DST      15      NaN
1  (UTC-05:00) US/Eastern;DST      15      NaN
```

```

      :@computed_region_f5dn_yrer :@computed_region_yeji_bk3q \
0          14          2
1          14          2

      :@computed_region_92fq_4b7q :@computed_region_sbqj_enih      counter
0          27          50  Y2H13094304
1          27          50  Y2H13094304

```

```

[23]: # check if all columns are merged without Key column, 'site'
print(f"True if all columns are successfully merged: {count.shape[1] + site.
      ↳shape[1] - 1 == data.shape[1]}")
print(f"Zero if all sites are successfully merged: {data.name.isna().
      ↳sum()}")

```

```

True if all columns are successfully merged:    True
Zero if all sites are successfully merged:      0

```

```

[24]: # what are the columns in our dataframe?
data.columns

```

```

[24]: Index(['id1', 'counts', 'date', 'status', 'site', 'id', 'name', 'latitude',
            'longitude', 'domain', 'timezone', 'interval',
            ':@computed_region_efsh_h5xi', ':@computed_region_f5dn_yrer',
            ':@computed_region_yeji_bk3q', ':@computed_region_92fq_4b7q',
            ':@computed_region_sbqj_enih', 'counter'],
            dtype='object')

```

```

[25]: # select required columns

data = data[['date', 'site', 'name', 'latitude', 'longitude', 'counts']]
data.reset_index(drop = True, inplace = True)
data.head(2)

```

```

[25]:
      date          site          name  latitude \
0  2021-01-01T00:15:00.000  100009425  Prospect Park West  40.67128846
1  2021-01-01T00:30:00.000  100009425  Prospect Park West  40.67128846

      longitude counts
0  -73.97138165      11
1  -73.97138165       8

```

```

[26]: # create a dataframe to gather information about each column
data_profiling_df = pd.DataFrame(columns = ["column_name",
                                             "column_type",
                                             "unique_values",
                                             "duplicate_values",
                                             "null_values",

```



```
"non_null_values",
"percent_null"])
```

```
[27]: # loop through each column to add rows to the data_profiling_df dataframe
for column in data.columns:

    info_dict = {}

    try:
        info_dict["column_name"] = column
        info_dict["column_type"] = data[column].dtypes
        info_dict["unique_values"] = len(data[column].unique())
        info_dict["duplicate_values"] = (data[column].shape[0] - data[column].
↪isna().sum()) - len(data[column].unique())
        info_dict["null_values"] = data[column].isna().sum()
        info_dict["non_null_values"] = data[column].shape[0] - data[column].
↪isna().sum()
        info_dict["percent_null"] = round((data[column].isna().sum()) /
↪(data[column].shape[0]), 3)

    except:
        print(f"unable to read column: {column}")

    data_profiling_df = data_profiling_df.append(info_dict, ignore_index=True)

data_profiling_df.sort_values(by = ['unique_values', "non_null_values"],
                               ascending = [False, False],
                               inplace=True)
```

```
[28]: data_profiling_df
```

```
[28]:  column_name  column_type  unique_values  duplicate_values  null_values  \
0         date         object          43679           576815           0
5        counts         object           336           620024          134
1         site         object            15           620479           0
2         name         object            15           620479           0
3    latitude         object            13           620481           0
4   longitude         object            13           620481           0

    non_null_values  percent_null
0         620494           0.0
5         620360           0.0
1         620494           0.0
2         620494           0.0
3         620494           0.0
4         620494           0.0
```

### 1.3 Step 3: Data Cleansing (Initial)

1. drop unneeded columns
2. drop duplicate rows
3. check for outliers

```
[29]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 620494 entries, 0 to 620493
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        620494 non-null  object
 1   site        620494 non-null  object
 2   name        620494 non-null  object
 3   latitude    620494 non-null  object
 4   longitude   620494 non-null  object
 5   counts      620360 non-null  object
dtypes: object(6)
memory usage: 28.4+ MB
```

```
[30]: data[data.duplicated()]
```

```
[30]: Empty DataFrame
Columns: [date, site, name, latitude, longitude, counts]
Index: []
```

```
[31]: # find number of duplicate rows

print(f"number of duplicate rows: {len(data[data.duplicated()])}")
```

```
number of duplicate rows: 0
```

```
[32]: # drop duplicate rows
      ## drop duplicates here
      ## print new shape of data

data = data.drop_duplicates(keep = 'first')
print(f"new shape of data: {data.shape}")
```

```
new shape of data: (620494, 6)
```

```
[33]: len(data[data.duplicated()])
```

```
[33]: 0
```

```
[34]: # find number of Null rows

print(data.isna().sum(), '\n')
```

```
print(f"shape of date: {data.shape}")
```

```
date          0
site          0
name          0
latitude      0
longitude     0
counts       134
dtype: int64
```

shape of date: (620494, 6)

```
[35]: # delete Null rows
```

```
data.dropna(inplace = True)
```

```
[36]: # print new shape of data
```

```
print(data.isna().sum(), '\n')
print(f"new shape of date: {data.shape}")
```

```
date          0
site          0
name          0
latitude      0
longitude     0
counts        0
dtype: int64
```

new shape of date: (620360, 6)

```
[37]: # data type change
```

```
df = pd.DataFrame()

df['date'] = data['date'].astype('datetime64[ns]').dt.date
df['time'] = data['date'].astype('datetime64[ns]').dt.time
df['site_address'] = data['name']
df['latitude'] = pd.to_numeric(data['latitude'])
df['longitude'] = pd.to_numeric(data['longitude'])
df['counts'] = data['counts'].astype(int)
```

```
[38]: data = df.copy()
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 620360 entries, 0 to 620493
Data columns (total 6 columns):
 #   Column          Non-Null Count  Dtype
---  -

```

```

0   date          620360 non-null object
1   time          620360 non-null object
2   site_address  620360 non-null object
3   latitude      620360 non-null float64
4   longitude     620360 non-null float64
5   counts        620360 non-null int32
dtypes: float64(2), int32(1), object(3)
memory usage: 30.8+ MB

```

```

[39]: # check the number of values having Zero
      # if True, None zero value
      data.all()

```

```

[39]: date          True
      time          True
      site_address  True
      latitude      True
      longitude     True
      counts        False
      dtype: bool

```

We do not delete these values to find the number of check-frequency per bicycle site

```

[40]: data.head(2)

```

```

[40]:      date      time      site_address  latitude  longitude  counts
0  2021-01-01  00:15:00  Prospect Park West  40.671288 -73.971382      11
1  2021-01-01  00:30:00  Prospect Park West  40.671288 -73.971382       8

```

## 1.4 Step 4: Additional data

### 1.4.1 Site Zipcode

```

[41]: # add Zipcode from external reference

      # get site names
      site_names = data.site_address.value_counts().reset_index()['index'].values
      print(f"Number of unique site name: {len(site_names)}")

```

Number of unique site name: 15

```

[42]: print(site_names[0:5], '\n')
      print(site_names[5:10], '\n')
      print(site_names[10:16], '\n')

```

```

['Prospect Park West' 'Manhattan Bridge Ped Path'
 'Williamsburg Bridge Bike Path' 'Ed Koch Queensboro Bridge Shared Path'
 'Staten Island Ferry']

```

```

['Pulaski Bridge' 'Kent Ave btw North 8th St and North 9th St']

```

```
'Brooklyn Bridge Bike Path' 'Manhattan Bridge Display Bike Counter'
'Amsterdam Ave at 86th St.']
```

```
['Columbus Ave at 86th St.' 'Manhattan Bridge Bike Comprehensive'
'Comprehensive Brooklyn Bridge Counter' '8th Ave at 50th St.'
'Brooklyn Bridge Bicycle Path (Roadway)']
```

Reference link: <https://www.unitedstateszipcodes.org/>

```
[43]: # To double-check the zipcode referencing with the site name, we do not use
      ↪ loop here.

data.loc[data['site_address'] == 'Prospect Park West', 'zipcode'] = 11215
data.loc[data['site_address'] == 'Manhattan Bridge Ped Path', 'zipcode'] = 10002
data.loc[data['site_address'] == 'Williamsburg Bridge Bike Path', 'zipcode'] =
      ↪ 10002
data.loc[data['site_address'] == 'Ed Koch Queensboro Bridge Shared Path',
      ↪ 'zipcode'] = 11101
data.loc[data['site_address'] == 'Staten Island Ferry', 'zipcode'] = 10301

data.loc[data['site_address'] == 'Pulaski Bridge', 'zipcode'] = 11101
data.loc[data['site_address'] == 'Kent Ave btw North 8th St and North 9th St',
      ↪ 'zipcode'] = 11249
data.loc[data['site_address'] == 'Brooklyn Bridge Bike Path', 'zipcode'] = 11214
data.loc[data['site_address'] == 'Manhattan Bridge Display Bike Counter',
      ↪ 'zipcode'] = 10002
data.loc[data['site_address'] == 'Amsterdam Ave at 86th St.', 'zipcode'] = 10024

data.loc[data['site_address'] == 'Columbus Ave at 86th St.', 'zipcode'] = 10024
data.loc[data['site_address'] == 'Manhattan Bridge Bike Comprehensive',
      ↪ 'zipcode'] = 10002
data.loc[data['site_address'] == 'Comprehensive Brooklyn Bridge Counter',
      ↪ 'zipcode'] = 10038
data.loc[data['site_address'] == '8th Ave at 50th St.', 'zipcode'] = 11220
data.loc[data['site_address'] == 'Brooklyn Bridge Bicycle Path (Roadway)',
      ↪ 'zipcode'] = 10038

# data type change
data['zipcode'] = data['zipcode'].astype(int)
data['zipcode'] = data['zipcode'].astype(str) # to merge with 'Zip per
      ↪ population' dataset
```

```
[44]: data.head(2)
```

```
[44]:      date      time  site_address  latitude  longitude  counts  \
0  2021-01-01  00:15:00  Prospect Park West  40.671288 -73.971382      11
```

```
1 2021-01-01 00:30:00 Prospect Park West 40.671288 -73.971382 8
```

```
    zipcode
0    11215
1    11215
```

#### 1.4.2 Population per zipcode

```
[45]: zip_pop = pd.read_csv('https://raw.githubusercontent.com/cpasean/Projects/main/
    ↪zip%20per%20population.csv')
zip_pop.head(2)
```

```
[45]: Rank Zip Code Population
0    1    11368    112,088
1    2    11385    107,796
```

```
[46]: # ignore warning message
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)
warnings.filterwarnings("ignore") # ignore runtime warning

# get only wanted columns
zip_pop = zip_pop[['Zip Code', 'Population']]
zip_pop.rename(columns = {'Zip Code' : 'zipcode', 'Population' : 'population'},
    ↪inplace = True)
```

```
[47]: data.shape
```

```
[47]: (620360, 7)
```

```
[48]: data = pd.merge(data, zip_pop, how = 'inner', on = 'zipcode')
data.head(2)
```

```
[48]:      date      time      site_address  latitude  longitude  counts  \
0  2021-01-01  00:15:00  Prospect Park West  40.671288 -73.971382     11
1  2021-01-01  00:30:00  Prospect Park West  40.671288 -73.971382      8

    zipcode  population
0    11215      69,873
1    11215      69,873
```

```
[49]: data.shape
```

```
[49]: (576689, 8)
```

## 1.5 Step 5: Data Cleansing (Final)

```
[50]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 576689 entries, 0 to 576688
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date            576689 non-null object
1   time            576689 non-null object
2   site_address    576689 non-null object
3   latitude        576689 non-null float64
4   longitude        576689 non-null float64
5   counts          576689 non-null int32
6   zipcode         576689 non-null object
7   population      576689 non-null object
dtypes: float64(2), int32(1), object(5)
memory usage: 37.4+ MB
```

```
[51]: # data type change as integer for zipcode and population
```

```
data['zipcode'] = data['zipcode'].astype(int)

pop = []
for i in data['population']:
    pop.append(int(i.replace(',', '')))
data['population'] = pop
```

```
[52]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 576689 entries, 0 to 576688
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date            576689 non-null object
1   time            576689 non-null object
2   site_address    576689 non-null object
3   latitude        576689 non-null float64
4   longitude        576689 non-null float64
5   counts          576689 non-null int32
6   zipcode         576689 non-null int32
7   population      576689 non-null int64
dtypes: float64(2), int32(2), int64(1), object(3)
memory usage: 35.2+ MB
```

```
[53]: # find number of duplicate rows

print(f"number of duplicate rows: {len(data[data.duplicated()])}")
```

number of duplicate rows: 0

```
[54]: # find number of Null rows

print(data.isna().sum(), '\n')
print(f"Final shape of date: {data.shape}")
```

```
date          0
time          0
site_address  0
latitude      0
longitude     0
counts        0
zipcode       0
population    0
dtype: int64
```

Final shape of date: (576689, 8)

```
[55]: data.head(2)
```

```
[55]:
```

	date	time	site_address	latitude	longitude	counts	\
0	2021-01-01	00:15:00	Prospect Park West	40.671288	-73.971382	11	
1	2021-01-01	00:30:00	Prospect Park West	40.671288	-73.971382	8	

	zipcode	population
0	11215	69873
1	11215	69873

## 1.6 Step 6: Create Location Dimension

```
[56]: # first, copy the entire table
```

```
location_dim = data.copy()
```

```
[57]: # second, subset for only the wanted columns in the dimension
```

```
location_dim = location_dim[['longitude',
                             'latitude',
                             'site_address',
                             'zipcode']]
```

```
[58]: # third, drop duplicate rows in dimension
```



```
location_dim = location_dim.drop_duplicates(subset = ["site_address"], keep = "first")
location_dim = location_dim.reset_index(drop = True)

location_dim.head()
```

```
[58]:
```

	longitude	latitude	site_address	zipcode
0	-73.971382	40.671288	Prospect Park West	11215
1	-73.994950	40.714573	Manhattan Bridge Ped Path	10002
2	-73.961450	40.710530	Williamsburg Bridge Bike Path	10002
3	-73.994750	40.715600	Manhattan Bridge Display Bike Counter	10002
4	-73.994750	40.715600	Manhattan Bridge Bike Comprehensive	10002

```
[59]: # fourth, add location_id as a surrogate key

location_dim.insert(0, "location_id",
                    range(10, 10+len(location_dim)))
```

```
[60]: location_dim.head(2)
```

```
[60]:
```

	location_id	longitude	latitude	site_address	zipcode
0	10	-73.971382	40.671288	Prospect Park West	11215
1	11	-73.994950	40.714573	Manhattan Bridge Ped Path	10002

```
[61]: # fifth, add the location_id to the Fact table

data = data.merge(location_dim[["site_address", "location_id"]],
                  left_on = "site_address",
                  right_on = "site_address",
                  how = "left")

data.head(2)

## merge with data to put location_id into Fact table
```

```
[61]:
```

	date	time	site_address	latitude	longitude	counts	\
0	2021-01-01	00:15:00	Prospect Park West	40.671288	-73.971382	11	
1	2021-01-01	00:30:00	Prospect Park West	40.671288	-73.971382	8	

	zipcode	population	location_id
0	11215	69873	10
1	11215	69873	10

```
[62]: # check latitude / longitude
```

```
import folium
```

```

lat = location_dim.latitude.mean()
long = location_dim.longitude.mean()

m = folium.Map(location=[lat, long], zoom_start = 9) # center location

for i in location_dim.index[:]:
    tooltip = location_dim.loc[i, "site_address"]
    lat = location_dim.loc[i, "latitude"]
    long = location_dim.loc[i, "longitude"]

    folium.Marker([lat, long], tooltip = tooltip).add_to(m)

# visual images are separately attached
m

```

[62]: <folium.folium.Map at 0x1cb32dfceb0>

## 1.7 Step 7: Create Time Dimension

[63]: # first, copy the entire table

```
time_dim = data.copy()
```

[64]: # second, subset for only the wanted columns in the dimension

```
time_dim = time_dim[['time']]
```

[65]: # third, drop duplicate rows in dimension

```

time_dim = time_dim.drop_duplicates(subset = ["time"], keep = "first")
time_dim = time_dim.reset_index(drop = True)

time_dim.head()

```

[65]:

	time
0	00:15:00
1	00:30:00
2	00:45:00
3	01:00:00
4	01:15:00

[66]: # fourth, add location\_id as a surrogate key

```
time_dim.insert(0, "time_id", range(100, 100+len(time_dim)))
```

[67]: time\_dim.tail(2)

```
[67]:      time_id      time
      94      194  23:45:00
      95      195   00:00:00
```

```
[68]: # fifth, add the location_id to the Fact table

data = data.merge(time_dim,
                  left_on = "time",
                  right_on = "time",
                  how = "left")

data.head(2)

## merge with data to put location_id into Fact table
```

```
[68]:      date      time      site_address  latitude  longitude  counts  \
0  2021-01-01  00:15:00  Prospect Park West  40.671288  -73.971382      11
1  2021-01-01  00:30:00  Prospect Park West  40.671288  -73.971382       8

      zipcode  population  location_id  time_id
0      11215      69873           10      100
1      11215      69873           10      101
```

## 1.8 Step 8: Create Date Dimension

```
[69]: # first, create a BigQuery client to connect to BigQuery
from google.cloud import bigquery
from google.oauth2 import service_account

key_path = r'C:\Users\aicpa\Google_
↳Drive\_CPADataScientistValueInvestor\_DataWarehouse\cis9440-340819-fdb3569fc29a.
↳json' # must edit to your credentials json file location
credentials = service_account.Credentials.from_service_account_file(key_path,
                                                                      ↳
↳scopes=["https://www.googleapis.com/auth/cloud-platform"],)
client = bigquery.Client(credentials = credentials,
                        project = credentials.project_id)
```

```
[70]: print(client)
```

```
<google.cloud.bigquery.client.Client object at 0x000001CB3B2A5C10>
```

```
[71]: sql_query = """
      SELECT
      CONCAT_
↳(FORMAT_DATE("%Y",d),FORMAT_DATE("%m",d),FORMAT_DATE("%d",d)) as date_id,
      d AS full_date,
      FORMAT_DATE('%w', d) AS week_day,
```

```

        FORMAT_DATE('%A', d) AS day_name,
        EXTRACT(DAY FROM d) AS year_day,
        EXTRACT(WEEK FROM d) AS week,
        EXTRACT(WEEK FROM d) AS year_week,
        EXTRACT(MONTH FROM d) AS month,
        FORMAT_DATE('%B', d) as month_name,
        FORMAT_DATE('%Q', d) as fiscal_qtr,
        EXTRACT(YEAR FROM d) AS year,
        (CASE WHEN FORMAT_DATE('%A', d) IN ('Sunday', 'Saturday') THEN 0
↪ELSE 1 END) AS day_is_weekday,
    FROM (
        SELECT
            *
        FROM
            UNNEST(GENERATE_DATE_ARRAY('2021-01-01', '2023-01-01', INTERVAL_
↪1 DAY)) AS d )
    """

# store extracted data in new dataframe
date_dim = client.query(sql_query).to_dataframe()

# validate that >0 stories have been extracted and return dataframe
if len(date_dim) > 0:
    print("date dimension created")
else:
    print("date dimension FAILED")

```

date dimension created

```
[72]: date_dim.head()
```

```
[72]:
```

	date_id	full_date	week_day	day_name	year_day	week	year_week	month	\
0	20210101	2021-01-01	5	Friday	1	0	0	1	
1	20210102	2021-01-02	6	Saturday	2	0	0	1	
2	20210103	2021-01-03	0	Sunday	3	1	1	1	
3	20210104	2021-01-04	1	Monday	4	1	1	1	
4	20210105	2021-01-05	2	Tuesday	5	1	1	1	

	month_name	fiscal_qtr	year	day_is_weekday
0	January	1	2021	1
1	January	1	2021	0
2	January	1	2021	0
3	January	1	2021	1
4	January	1	2021	1

```
[73]: # create date_id column in the Fact Table
```

```
data['date_id'] = data['date'].apply(lambda x: pd.to_datetime(x).
↳strftime("%Y%m%d"))
```

## 1.9 Step 9: Creating Fact(s)

```
[74]: # Creating Bicycle Fact Table

fact_bicycle = data[["date_id",
                    "time_id",
                    "location_id",
                    "counts"]]
fact_bicycle.sample(5)
```

```
[74]:      date_id  time_id  location_id  counts
175605  20210110      157           14       81
372684  20210831      188           18       39
298200  20220112      182           16       19
283296  20210810      158           16       20
280446  20210711      188           16       23
```

```
[75]: # Creating Census Fact Table

fact_census = data[["date_id",
                    "location_id",
                    "population"]]
fact_census.sample(5)
```

```
[75]:      date_id  location_id  population
163750  20211208           13       74479
515792  20210119           22       23311
205720  20211120           14       74479
226197  20210323           15       31366
544155  20211111           22       23311
```

## 1.10 Step 10: Deliver Facts and Dimensions to Data Warehouse (BigQuery)

```
[76]: # build a function to load tables to BigQuery

def load_table_to_bigquery(df, table_name):

    dataset_id = 'cis9440-340819.final_project_etl_nyc_bicycle'

    dataset_ref = client.dataset(dataset_id)
    job_config = bigquery.LoadJobConfig()
    job_config.autodetect = True
    job_config.write_disposition = "WRITE_TRUNCATE"
```

```

upload_table_name = f"cis9440-340819.final_project_etl_nyc_bicycle.
↳{table_name}"

load_job = client.load_table_from_dataframe(df,
                                           upload_table_name,
                                           job_config = job_config)

print(f"starting job {load_job}")

```

```

[77]: load_table_to_bigquery(df = location_dim,
                             table_name = "location_dim")

```

starting job LoadJob<project=cis9440-340819, location=US, id=40b6c24f-3550-41ad-bcdb-6e044b5a986f>

```

[78]: load_table_to_bigquery(df = time_dim,
                             table_name = "time_dim")

```

starting job LoadJob<project=cis9440-340819, location=US, id=c1c79eef-5419-48d6-9886-8c373f7eff59>

```

[79]: load_table_to_bigquery(df = date_dim,
                             table_name = "date_dim")

```

starting job LoadJob<project=cis9440-340819, location=US, id=b63f7c91-7a19-42d7-b34d-129977c7318a>

```

[80]: load_table_to_bigquery(df = fact_bicycle,
                             table_name = "fact_bicycle")

```

starting job LoadJob<project=cis9440-340819, location=US, id=9072e453-baba-42d2-834a-9cff73eaf8b3>

```

[81]: load_table_to_bigquery(df = fact_census,
                             table_name = "fact_census")

```

starting job LoadJob<project=cis9440-340819, location=US, id=a820ab20-2460-494c-88eb-ec465bb7ad06>

```

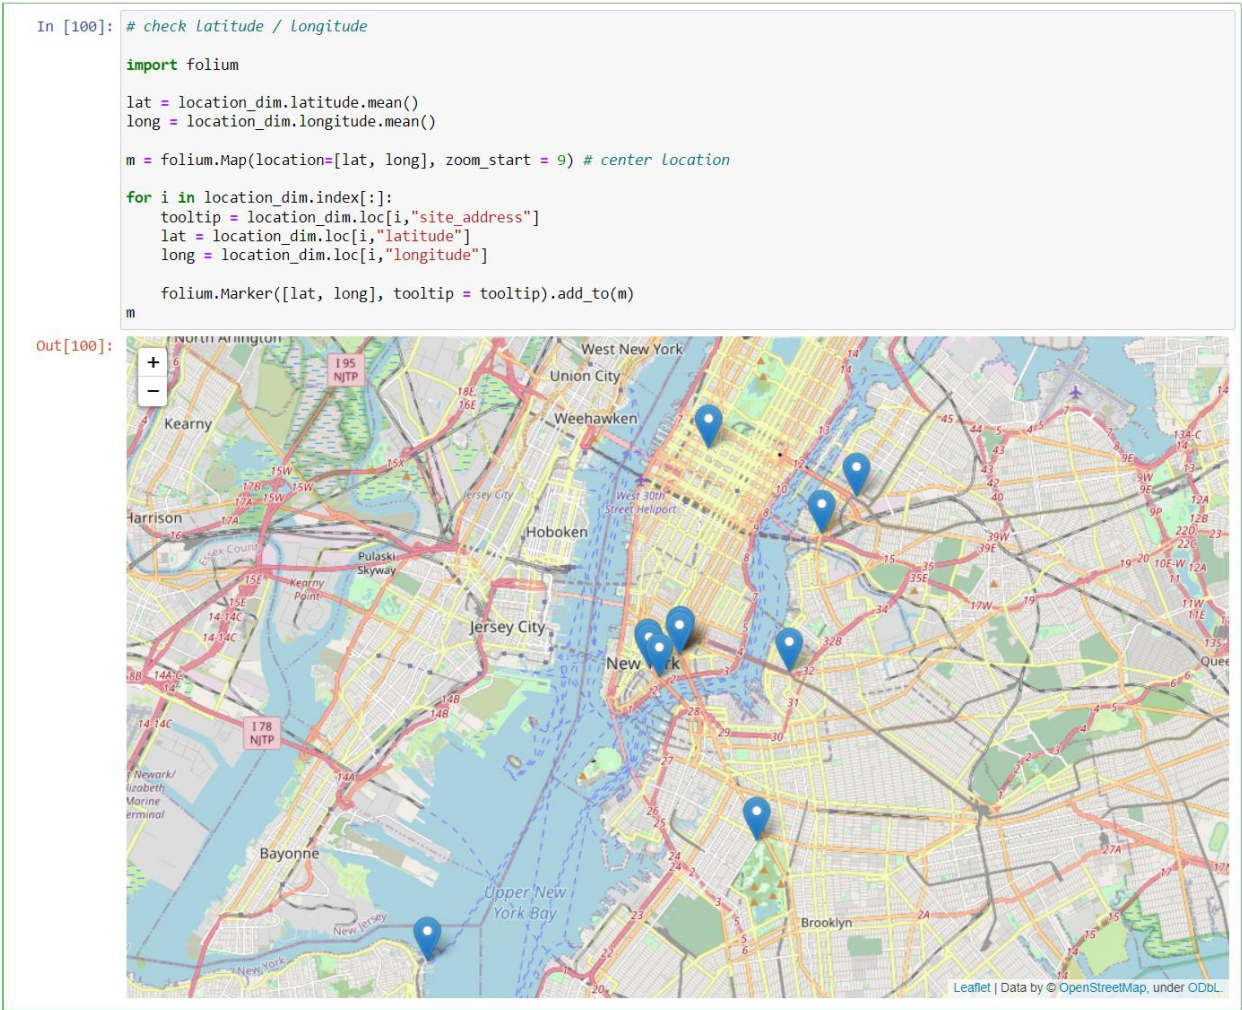
[82]: # send to *.csv

location_dim.to_csv('location_dim.csv')
time_dim.to_csv('time_dim.csv')
date_dim.to_csv('date_dim.csv')
fact_bicycle.to_csv('fact_bicycle.csv')
fact_census.to_csv('fact_census.csv')

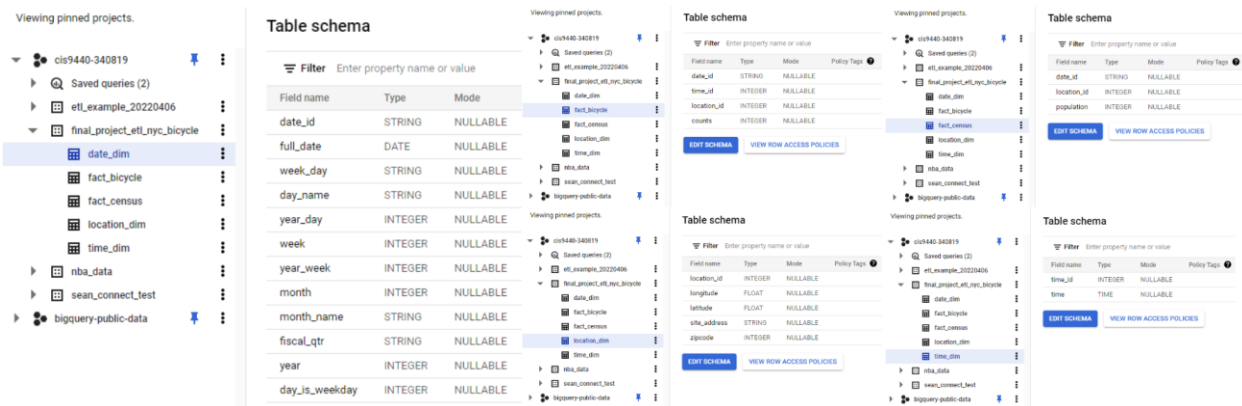
```

## 1.11 Appendix

Test Image #1: Site location




Test Image #2: Big Query



### Test Image #3: Tableau

fact\_bicycle+ (final\_project\_etl\_nyc\_bicycle)



fact\_census+ (final\_project\_etl\_nyc\_bicycle)

